

JD-specific resume optimization via coding agent

Same experience bank, a different resume for each posting.

github.com/henryzhongsc/comp608_resume_via_agent

Why optimize a resume against a job description

- A job description often describes the real role poorly.
- HR teams and automated screening systems still rely on it heavily.
- Screening is increasingly AI-powered, so the wording of a resume matters more every year.
- A qualified candidate can be filtered out over wording rather than over ability.
- Early-career applicants lose the most, because they have the least track record to compensate with.

Why use a coding agent

- Commercial services such as Jobright already optimize a resume in one click.
- One click produces one output, with no control over it and no room to refine it.
- A coding agent can revise the same resume repeatedly until it is right.
- The differences between versions show how you read each job, and they are worth keeping.
- A git repo records every version, while a third-party service does not.
- Most of you already pay for a coding plan.

You need a detailed experience bank first

- An agent can only use experience that you have written down somewhere.
- The pre-class assignment exists to build that written record.
- Many resumes over-focus on the numbers and tech stack the writer wants to show off.
- That focus loses the problem being solved, the scope that was owned, and what the work shows about the person.
- Recording everything in detail lets the agent decide what belongs on the page.

What to do today

- Objective 1: run the loop on Fred's material.
 - Clone the repo and read the readme:
`github.com/henryzhongsc/comp608_resume_via_agent` — set it to private if you wish.
 - Add a job description you are actually applying to under `jd/`.
 - Fire up a coding agent such as Claude Code, run `/optimize-against-jd` on that posting, and compare the result against the other resumes.
- Objective 2: replace Fred with yourself.
 - Replace the identity, experiences, and projects under `raw_info/` with your own, and define the constants for each.
 - Run the skill again to get a real resume. Unfinished is fine.
 - Judge whether the result is what you want. If it is not, iterate with the agent, and fold what you learn back into the `/optimize-against-jd` skill.

Why a skill, and why constants

- Targeting a resume is a repeated task, so the instructions belong in a reusable skill.
 - Every run is a chance to improve the instructions in that skill.
- Constants declare what must appear no matter which job is targeted.
 - This repo declares them per experience and per project.
 - They can also decide which experiences and projects appear at all.
- Several constant sets can be prepared and passed to the skill as an input.